

04 - Allen-Cahn Equation Experiments

AC-PINN Project | Authors: Suyash Vasal Jain, Nishita Raghvendra

PDE: $u_t = \epsilon^2 u_{xx} + u - u^3$, $\epsilon = 0.01$

IC: $u(x,0) = x^2 \cos(\pi x)$ | **BC:** $u(-1,t) = u(1,t) = 0$

Domain: $x \in [-1,1]$, $t \in [0,1]$

Architecture: [2, 128, 128, 128, 128, 128, 1] | **Epochs:** 10000

Note: Allen-Cahn is the hardest test. Stiff PDE with sharp interface. Vanilla PINN is known to fail here. This is AC-PINN's strongest showcase.

```
In [1]: import sys, os
sys.path.append('.')
import torch
import numpy as np
import matplotlib.pyplot as plt
from pinn_base import (
    device, NoisyDataGenerator, PINNSolver, ACPINNSolver,
    AllenCahnFDM, Benchmark, save_metrics, save_history, save_training_plots
)

PDE = 'allen_cahn'
LAYERS = [2, 128, 128, 128, 128, 128, 1]
EPOCHS = 10000
PDE_PARAMS = {'epsilon': 0.01}
RESULTS = '../results/allen_cahn/'
FIGURES = '../figures/allen_cahn/'
os.makedirs(RESULTS, exist_ok=True)
os.makedirs(FIGURES, exist_ok=True)


gen = NoisyDataGenerator(pde=PDE)
print(f'Device: {device}')
```

Device: cuda

Section 1 - FDM Ground Truth

```
In [2]: fdm = AllenCahnFDM(nx=256, nt=5000, epsilon=0.01)
fdm.solve()
fdm.plot_solution(title='Allen-Cahn FDM - Ground Truth')
fdm.plot_time_slices()
```

AllenCahnFDM solved in 1.9181s



Section 2 - Data Conditions

```
In [3]: # Allen-Cahn needs more collocation points due to stiffness
data_clean_dense = gen.generate(N_ic=1000, N_bc=1000, N_f=10000, noise_eps=0.0)
data_noisy_sparse = gen.generate(N_ic=20, N_bc=20, N_f=3000, noise_eps=0.1)
print('Data ready')
```

Data ready


Section 3 - Vanilla PINN, Clean Dense

```
In [4]: # Higher lambda_pde for stiff PDE
vanilla_clean = PINNSolver(pde=PDE, layers=LAYERS, pde_params=PDE_PARAMS,
                           lambda_ic=1.0, lambda_bc=1.0, lambda_pde=10.0)
h_vc = vanilla_clean.fit(data_clean_dense, epochs=EPOCHS, print_every=1000, label='
save_training_plots(h_vc, FIGURES+'vanilla_clean_training.png', 'Allen-Cahn | Vanil
vanilla_clean.plot_loss_history(h_vc)
vanilla_clean.plot_solution(title='Vanilla PINN - Allen-Cahn Clean Dense')
vanilla_clean.plot_initial_condition_comparison(gen)
save_history(h_vc, RESULTS+'vanilla_clean_history.npy')
```

D:\PINN\ac-pinn-project\ac-pinn-project\venv\Lib\site-packages\torch\autograd\graph.py:825: UserWarning: Attempting to run cuBLAS, but there was no current CUDA context! Attempting to set the primary context... (Triggered internally at C:\actions-runner_work\pytorch\pytorch\builder\windows\pytorch\aten\src\ATen\cuda\CublasHandlePool.cpp:135.)

return Variable._execution_engine.run_backward(# Calls into the C++ engine to run the backward pass

```
[Allen-Cahn | Vanilla Clean | Epoch    0/10000] Total: 0.626974 | IC: 0.144497 | B
C: 0.087670 | PDE: 0.039481
[Allen-Cahn | Vanilla Clean | Epoch  1000/10000] Total: 0.128943 | IC: 0.121899 | B
C: 0.003938 | PDE: 0.000311
[Allen-Cahn | Vanilla Clean | Epoch  2000/10000] Total: 0.130245 | IC: 0.124084 | B
C: 0.004663 | PDE: 0.000150
[Allen-Cahn | Vanilla Clean | Epoch  3000/10000] Total: 0.128454 | IC: 0.121195 | B
C: 0.003796 | PDE: 0.000346
[Allen-Cahn | Vanilla Clean | Epoch  4000/10000] Total: 0.124353 | IC: 0.117401 | B
C: 0.001691 | PDE: 0.000526
[Allen-Cahn | Vanilla Clean | Epoch  5000/10000] Total: 0.074803 | IC: 0.068328 | B
C: 0.002227 | PDE: 0.000425
[Allen-Cahn | Vanilla Clean | Epoch  6000/10000] Total: 0.048437 | IC: 0.040147 | B
C: 0.003748 | PDE: 0.000454
[Allen-Cahn | Vanilla Clean | Epoch  7000/10000] Total: 0.034409 | IC: 0.023831 | B
C: 0.006994 | PDE: 0.000358
[Allen-Cahn | Vanilla Clean | Epoch  8000/10000] Total: 0.019724 | IC: 0.017373 | B
C: 0.000474 | PDE: 0.000188
[Allen-Cahn | Vanilla Clean | Epoch  9000/10000] Total: 0.020294 | IC: 0.015079 | B
C: 0.003662 | PDE: 0.000155
Training complete in 276.59s
```





N/A - Vanilla PINN

N/A - Vanilla PINN

Saved: ../figures/allen_cahn/vanilla_clean_training.png

Saved: ../figures/allen_cahn/vanilla_clean_training.csv





Saved: ../results/allen_cahn/vanilla_clean_history.npy

Section 4 - Vanilla PINN, Noisy Sparse


⚠ Vanilla PINN typically fails here on Allen-Cahn.

```
In [5]: vanilla_noisy = PINNSolver(pde=PDE, layers=LAYERS, pde_params=PDE_PARAMS,
                                   lambda_ic=1.0, lambda_bc=1.0, lambda_pde=10.0)
h_vn = vanilla_noisy.fit(data_noisy_sparse, epochs=EPOCHS, print_every=1000, label=
save_training_plots(h_vn, FIGURES+'vanilla_noisy_training.png', 'Allen-Cahn | Vanil
vanilla_noisy.plot_loss_history(h_vn)
vanilla_noisy.plot_solution(title='Vanilla PINN - Allen-Cahn Noisy Sparse')
save_history(h_vn, RESULTS+'vanilla_noisy_history.npy')
```

```



[Allen-Cahn | Vanilla Noisy | Epoch    0/10000] Total: 0.222897 | IC: 0.160220 | B
C: 0.031515 | PDE: 0.003116
[Allen-Cahn | Vanilla Noisy | Epoch  1000/10000] Total: 0.171356 | IC: 0.147896 | B
C: 0.018931 | PDE: 0.000453
[Allen-Cahn | Vanilla Noisy | Epoch  2000/10000] Total: 0.175861 | IC: 0.149349 | B
C: 0.023983 | PDE: 0.000253
[Allen-Cahn | Vanilla Noisy | Epoch  3000/10000] Total: 0.143990 | IC: 0.113653 | B
C: 0.025421 | PDE: 0.000492
[Allen-Cahn | Vanilla Noisy | Epoch  4000/10000] Total: 0.130944 | IC: 0.102572 | B
C: 0.020978 | PDE: 0.000739
[Allen-Cahn | Vanilla Noisy | Epoch  5000/10000] Total: 0.120333 | IC: 0.096492 | B
C: 0.017651 | PDE: 0.000619
[Allen-Cahn | Vanilla Noisy | Epoch  6000/10000] Total: 0.127031 | IC: 0.105304 | B
C: 0.017335 | PDE: 0.000439
[Allen-Cahn | Vanilla Noisy | Epoch  7000/10000] Total: 0.102742 | IC: 0.075615 | B
C: 0.019211 | PDE: 0.000792
[Allen-Cahn | Vanilla Noisy | Epoch  8000/10000] Total: 0.059195 | IC: 0.033552 | B
C: 0.016160 | PDE: 0.000948
[Allen-Cahn | Vanilla Noisy | Epoch  9000/10000] Total: 0.025510 | IC: 0.008121 | B
C: 0.014946 | PDE: 0.000244
Training complete in 136.67s

```



Saved: ../figures/allen_cahn/vanilla_noisy_training.png

Saved: ../figures/allen_cahn/vanilla_noisy_training.csv



Saved: ../results/allen_cahn/vanilla_noisy_history.npy


Section 5 - AC-PINN, Clean Dense

```
In [6]: ac_clean = ACPINNSolver(pde=PDE, layers=LAYERS, pde_params=PDE_PARAMS,
                                weight_strategy='both',
                                N_pool=30000, resample_every=500)
h_ac = ac_clean.fit(data_clean_dense, epochs=EPOCHS, print_every=1000, label='Allen
save_training_plots(h_ac, FIGURES+'acpinn_clean_training.png', 'Allen-Cahn | AC-PIN
ac_clean.plot_loss_history(h_ac)
ac_clean.plot_weight_history(h_ac)
ac_clean.plot_solution(title='AC-PINN - Allen-Cahn Clean Dense')
```

```
ac_clean.plot_initial_condition_comparison(gen)
save_history(h_ac, RESULTS+'ac_clean_history.npy')
```




[Allen-Cahn | AC-PINN Clean | Epoch 0/10000] Stage 1/4 | Total: 0.196306 | IC: 0.142633 | BC: 0.049597 | PDE: 0.000815 | $\lambda=(1.00,1.00,5.00)$
 [Allen-Cahn | AC-PINN Clean | Epoch 1000/10000] Stage 1/4 | Total: 0.091380 | IC: 0.006492 | BC: 0.029534 | PDE: 0.002163 | $\lambda=(4.79,1.59,6.19)$
 [Allen-Cahn | AC-PINN Clean | Epoch 2000/10000] Stage 1/4 | Total: 0.013632 | IC: 0.000618 | BC: 0.006259 | PDE: 0.000002 | $\lambda=(2.38,1.94,10.00)$
 [Allen-Cahn | AC-PINN Clean | Epoch 3000/10000] Stage 2/4 | Total: 0.009599 | IC: 0.000539 | BC: 0.005325 | PDE: 0.000000 | $\lambda=(3.57,1.44,10.00)$
 [Allen-Cahn | AC-PINN Clean | Epoch 4000/10000] Stage 2/4 | Total: 0.007876 | IC: 0.000050 | BC: 0.006308 | PDE: 0.000000 | $\lambda=(6.56,1.20,10.00)$
 [Allen-Cahn | AC-PINN Clean | Epoch 5000/10000] Stage 2/4 | Total: 0.012419 | IC: 0.000160 | BC: 0.009579 | PDE: 0.000000 | $\lambda=(9.12,1.14,10.00)$
 [Allen-Cahn | AC-PINN Clean | Epoch 6000/10000] Stage 3/4 | Total: 0.023954 | IC: 0.000041 | BC: 0.008028 | PDE: 0.000004 | $\lambda=(0.10,2.98,0.10)$
 [Allen-Cahn | AC-PINN Clean | Epoch 7000/10000] Stage 3/4 | Total: 0.143955 | IC: 0.049044 | BC: 0.001121 | PDE: 0.000001 | $\lambda=(2.93,0.10,0.10)$
 [Allen-Cahn | AC-PINN Clean | Epoch 8000/10000] Stage 4/4 | Total: 0.175603 | IC: 0.050622 | BC: 0.067296 | PDE: 0.055103 | $\lambda=(0.88,1.17,0.96)$
 [Allen-Cahn | AC-PINN Clean | Epoch 9000/10000] Stage 4/4 | Total: 0.397028 | IC: 0.132734 | BC: 0.000153 | PDE: 0.000254 | $\lambda=(2.99,0.10,0.10)$
 AC-PINN training complete in 299.13s


Allen-Cahn | AC-PINN Clean



Saved: ../figures/allen_cahn/acpinn_clean_training.png

Saved: ../figures/allen_cahn/acpinn_clean_training.csv





Saved: ../results/allen_cahn/ac_clean_history.npy


Section 6 - AC-PINN, Noisy Sparse

This is where AC-PINN should demonstrate its biggest advantage.

```
In [7]: ac_noisy = ACPINNSolver(pde=PDE, layers=LAYERS, pde_params=PDE_PARAMS,
                                weight_strategy='both',
                                N_pool=30000, resample_every=500)
h_an = ac_noisy.fit(data_noisy_sparse, epochs=EPOCHS, print_every=1000, label='Allen-Cahn | AC-PINN')
save_training_plots(h_an, FIGURES+'acpinn_noisy_training.png', 'Allen-Cahn | AC-PINN')
ac_noisy.plot_loss_history(h_an)
ac_noisy.plot_weight_history(h_an)
ac_noisy.plot_solution(title='AC-PINN - Allen-Cahn Noisy Sparse')
save_history(h_an, RESULTS+'ac_noisy_history.npy')
```

```
[Allen-Cahn | AC-PINN Noisy | Epoch    0/10000] Stage 1/4 | Total: 0.211765 | IC:
0.150837 | BC: 0.058899 | PDE: 0.000406 |  $\lambda=(1.00,1.00,5.00)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 1000/10000] Stage 1/4 | Total: 0.039671 | IC:
0.007976 | BC: 0.009067 | PDE: 0.000087 |  $\lambda=(2.82,1.80,10.00)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 2000/10000] Stage 1/4 | Total: 0.039165 | IC:
0.006673 | BC: 0.007338 | PDE: 0.000016 |  $\lambda=(4.18,1.51,10.00)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 3000/10000] Stage 2/4 | Total: 0.026770 | IC:
0.005654 | BC: 0.006264 | PDE: 0.000068 |  $\lambda=(2.13,2.24,10.00)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 4000/10000] Stage 2/4 | Total: 0.030289 | IC:
0.005303 | BC: 0.006176 | PDE: 0.000036 |  $\lambda=(2.08,3.09,5.13)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 5000/10000] Stage 2/4 | Total: 0.025561 | IC:
0.004796 | BC: 0.006078 | PDE: 0.000023 |  $\lambda=(3.15,1.68,10.00)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 6000/10000] Stage 3/4 | Total: 0.017738 | IC:
0.005808 | BC: 0.006425 | PDE: 0.000495 |  $\lambda=(1.37,1.51,0.12)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 7000/10000] Stage 3/4 | Total: 0.016336 | IC:
0.005461 | BC: 0.006119 | PDE: 0.000933 |  $\lambda=(1.31,1.47,0.22)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 8000/10000] Stage 4/4 | Total: 1.036718 | IC:
0.019074 | BC: 0.032861 | PDE: 0.388097 |  $\lambda=(0.13,0.22,2.65)$ 
[Allen-Cahn | AC-PINN Noisy | Epoch 9000/10000] Stage 4/4 | Total: 0.441073 | IC:
0.159359 | BC: 0.014571 | PDE: 0.000252 |  $\lambda=(2.74,0.25,0.10)$ 
AC-PINN training complete in 158.41s
```


Allen-Cahn | AC-PINN Noisy




Saved: ../figures/allen_cahn/acpinn_noisy_training.png

Saved: ../figures/allen_cahn/acpinn_noisy_training.csv

Training Loss Curves





Saved: ../results/allen_cahn/ac_noisy_history.npy




Section 7 - Benchmark vs FDM

```
In [8]: bench = Benchmark(fdm, nx=200, nt=100)
bench.add('Vanilla (clean)', vanilla_clean)
bench.add('Vanilla (noisy)', vanilla_noisy)
bench.add('AC-PINN (clean)', ac_clean)
bench.add('AC-PINN (noisy)', ac_noisy)
bench.run()

metrics = bench.compare_metrics()
bench.plot_comparison(save_path=FIGURES+'comparison.png')
bench.plot_time_slices(save_path=FIGURES+'time_slices.png')
bench.plot_error_over_time(save_path=FIGURES+'error_over_time.png')

save_metrics(metrics, RESULTS+'benchmark_metrics.npy')
print('Allen-Cahn experiments complete.')
```

Model	Rel L2	Max Err	MAE	RMSE
Vanilla (clean)	0.143559	1.017695	0.035595	0.062932
Vanilla (noisy)	0.316406	1.047791	0.077888	0.138703
AC-PINN (clean)	0.812069	0.922218	0.229725	0.355987
AC-PINN (noisy)	1.040699	1.153148	0.309040	0.456213



Saved: ../results/allen_cahn/benchmark_metrics.npy
Allen-Cahn experiments complete.

Section 8 - Noise Level Study

```
In [9]: noise_results = {}
for eps in [0.05, 0.1, 0.2]:
    print(f'\n--- ε={eps} ---')
    d = gen.generate(N_ic=50, N_bc=50, N_f=5000, noise_eps=eps)
    v = PINNSolver(pde=PDE, layers=LAYERS, pde_params=PDE_PARAMS, lambda_pde=10.0)
```

```

v.fit(d, epochs=5000, print_every=2500)
a = ACPINNSolver(pde=PDE, layers=LAYERS, pde_params=PDE_PARAMS, weight_strategy='adam')
a.fit(d, epochs=5000, print_every=2500)
b = Benchmark(fdm).add(f'Vanilla  $\epsilon=\{\epsilon\}$ ', v).add(f'AC-PINN  $\epsilon=\{\epsilon\}$ ', a)
b.run()
noise_results[ $\epsilon$ ] = b.compare_metrics()
save_metrics(noise_results, RESULTS+'noise_study_metrics.npy')
print('Allen-Cahn noise study complete.')

```

--- $\epsilon=0.05$ ---

[Epoch 0/5000] Total: 0.591497 | IC: 0.160854 | BC: 0.155205 | PDE: 0.027544

[Epoch 2500/5000] Total: 0.124157 | IC: 0.112330 | BC: 0.009061 | PDE: 0.000277

Training complete in 86.70s

[Epoch 0/5000] Stage 1/4 | Total: 0.127733 | IC: 0.119624 | BC: 0.007746 | PDE: 0.000073 | $\lambda=(1.00, 1.00, 5.00)$

[Epoch 2500/5000] Stage 2/4 | Total: 0.015569 | IC: 0.002428 | BC: 0.002987 | PDE: 0.000005 | $\lambda=(3.77, 2.14, 3.74)$

AC-PINN training complete in 93.99s

Model	Rel L2	Max Err	MAE	RMSE
Vanilla $\epsilon=0.05$	0.549171	1.018321	0.132007	0.240741
AC-PINN $\epsilon=0.05$	0.848223	0.958067	0.241095	0.371837

--- $\epsilon=0.1$ ---

[Epoch 0/5000] Total: 0.811225 | IC: 0.213705 | BC: 0.222257 | PDE: 0.037526

[Epoch 2500/5000] Total: 0.159354 | IC: 0.137152 | BC: 0.018529 | PDE: 0.000367

Training complete in 85.38s

[Epoch 0/5000] Stage 1/4 | Total: 0.185915 | IC: 0.159454 | BC: 0.025178 | PDE: 0.000256 | $\lambda=(1.00, 1.00, 5.00)$

[Epoch 2500/5000] Stage 2/4 | Total: 0.045237 | IC: 0.005578 | BC: 0.012597 | PDE: 0.000016 | $\lambda=(4.60, 1.54, 7.43)$

AC-PINN training complete in 88.18s

Model	Rel L2	Max Err	MAE	RMSE
Vanilla $\epsilon=0.1$	0.794291	0.974870	0.248946	0.348194
AC-PINN $\epsilon=0.1$	0.979001	0.990372	0.288530	0.429166

--- $\epsilon=0.2$ ---

[Epoch 0/5000] Total: 0.268482 | IC: 0.125368 | BC: 0.076212 | PDE: 0.006690

[Epoch 2500/5000] Total: 0.178392 | IC: 0.086120 | BC: 0.069230 | PDE: 0.002304

Training complete in 81.85s

[Epoch 0/5000] Stage 1/4 | Total: 0.193024 | IC: 0.121321 | BC: 0.070732 | PDE: 0.000194 | $\lambda=(1.00, 1.00, 5.00)$

[Epoch 2500/5000] Stage 2/4 | Total: 0.233589 | IC: 0.042495 | BC: 0.063939 | PDE: 0.000048 | $\lambda=(2.78, 1.80, 10.00)$

AC-PINN training complete in 88.19s

Model	Rel L2	Max Err	MAE	RMSE
Vanilla $\epsilon=0.2$	0.459016	1.069122	0.109629	0.201219
AC-PINN $\epsilon=0.2$	0.770576	1.206568	0.217735	0.337798

Saved: ../results/allen_cahn/noise_study_metrics.npy

Allen-Cahn noise study complete.